

```
In [1]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.datasets import load_iris
iris = load_iris()

from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
from sklearn.svm import SVC

from sklearn.metrics import confusion_matrix
from sklearn.metrics import classification_report
```

dir(iris) — Inspecting Available Attributes and Methods

```
In [2]: dir(iris)
```

```
Out[2]: ['DESCR',
        'data',
        'data_module',
        'feature_names',
        'filename',
        'frame',
        'target',
        'target_names']
```

```
In [3]: iris.feature_names
```

```
Out[3]: ['sepal length (cm)',
        'sepal width (cm)',
        'petal length (cm)',
        'petal width (cm)']
```

```
In [4]: df = pd.DataFrame(iris.data, columns = iris.feature_names)
df
```

Out [4]:

	sepal length (cm)	sepal width (cm)	petal length (cm)	petal width (cm)
0	5.1	3.5	1.4	0.2
1	4.9	3.0	1.4	0.2
2	4.7	3.2	1.3	0.2
3	4.6	3.1	1.5	0.2
4	5.0	3.6	1.4	0.2
...
145	6.7	3.0	5.2	2.3
146	6.3	2.5	5.0	1.9
147	6.5	3.0	5.2	2.0
148	6.2	3.4	5.4	2.3
149	5.9	3.0	5.1	1.8

150 rows x 4 columns

In [5]: `df['target'] = iris.target`
`df`

Out [5]:

	sepal length (cm)	sepal width (cm)	petal length (cm)	petal width (cm)	target
0	5.1	3.5	1.4	0.2	0
1	4.9	3.0	1.4	0.2	0
2	4.7	3.2	1.3	0.2	0
3	4.6	3.1	1.5	0.2	0
4	5.0	3.6	1.4	0.2	0
...
145	6.7	3.0	5.2	2.3	2
146	6.3	2.5	5.0	1.9	2
147	6.5	3.0	5.2	2.0	2
148	6.2	3.4	5.4	2.3	2
149	5.9	3.0	5.1	1.8	2

150 rows x 5 columns

In [6]: `iris.target_names`

Out [6]: `array(['setosa', 'versicolor', 'virginica'], dtype='<U10')`

In [7]: `df['flower_name'] = df.target.apply(lambda x: iris.target_names[x])`

In [8]: `df`

Out [8]:

	sepal length (cm)	sepal width (cm)	petal length (cm)	petal width (cm)	target	flower_name
0	5.1	3.5	1.4	0.2	0	setosa
1	4.9	3.0	1.4	0.2	0	setosa
2	4.7	3.2	1.3	0.2	0	setosa
3	4.6	3.1	1.5	0.2	0	setosa
4	5.0	3.6	1.4	0.2	0	setosa
...
145	6.7	3.0	5.2	2.3	2	virginica
146	6.3	2.5	5.0	1.9	2	virginica
147	6.5	3.0	5.2	2.0	2	virginica
148	6.2	3.4	5.4	2.3	2	virginica
149	5.9	3.0	5.1	1.8	2	virginica

150 rows × 6 columns

Splitting the Dataset by Iris Species

```
In [9]: df0 = df[df.target==0]
df1 = df[df.target==1]
df2 = df[df.target==2]
```

In [10]: df0.head()

Out [10]:

	sepal length (cm)	sepal width (cm)	petal length (cm)	petal width (cm)	target	flower_name
0	5.1	3.5	1.4	0.2	0	setosa
1	4.9	3.0	1.4	0.2	0	setosa
2	4.7	3.2	1.3	0.2	0	setosa
3	4.6	3.1	1.5	0.2	0	setosa
4	5.0	3.6	1.4	0.2	0	setosa

In [11]: df1.head()

Out [11]:

	sepal length (cm)	sepal width (cm)	petal length (cm)	petal width (cm)	target	flower_name
50	7.0	3.2	4.7	1.4	1	versicolor
51	6.4	3.2	4.5	1.5	1	versicolor
52	6.9	3.1	4.9	1.5	1	versicolor
53	5.5	2.3	4.0	1.3	1	versicolor
54	6.5	2.8	4.6	1.5	1	versicolor

In [12]: `df2.head()`

Out[12]:


	sepal length (cm)	sepal width (cm)	petal length (cm)	petal width (cm)	target	flower_name
100	6.3	3.3	6.0	2.5	2	virginica
101	5.8	2.7	5.1	1.9	2	virginica
102	7.1	3.0	5.9	2.1	2	virginica
103	6.3	2.9	5.6	1.8	2	virginica
104	6.5	3.0	5.8	2.2	2	virginica

Scatter Plot of Sepal Length vs Sepal Width by Species

In [13]:

```
plt.scatter(df0['sepal length (cm)'],df0['sepal width (cm)'], color='green',marker='+')
plt.scatter(df1['sepal length (cm)'],df1['sepal width (cm)'], color='blue',marker='o')
plt.xlabel('sepal length (cm)')
plt.ylabel('sepal width (cm)')
```

Out[13]: Text(0, 0.5, 'sepal width (cm)')




Scatter Plot of petal Length vs petal Width by Species

In [14]:

```
plt.scatter(df0['petal length (cm)'],df0['petal width (cm)'],color='green',marker='+')
plt.scatter(df1['petal length (cm)'],df1['petal width (cm)'],color='blue',marker='o')
plt.xlabel('petal length (cm)')
plt.ylabel('petal width (cm)')
```

Out[14]: Text(0, 0.5, 'petal width (cm)')



🔧 Preparing Features and Labels for Machine Learning

```
In [15]: X = df.drop(['flower_name', 'target'], axis=1)
         y = df['target']
```

✂️ Splitting the Dataset into Training and Testing Sets

```
In [16]: X_train, X_test, y_train, y_test = train_test_split(X, y, train_size=0.8)
```

```
In [17]: len(X_train)
```

Out[17]: 120

```
In [18]: len(X_test)
```

Out[18]: 30

Training a Support Vector Machine (SVM) Classifier What SVM does:

It tries to find the best hyperplane that separates the classes in the feature space.

By default, it uses the RBF (Radial Basis Function) kernel, which can handle non-linear boundaries.

```
In [19]: model = SVC()
```

```
model.fit(X_train, y_train)
```

Out [19]:


```
▼ SVC  
SVC()
```

Evaluating the Model Accuracy on Test Data

```
In [20]: model.score(X_test, y_test)
```

Out [20]: 0.9666666666666667

Training SVM with Regularization Parameter C=10 SVC(C=10) creates a Support Vector Classifier with a regularization parameter C set to 10. .

 **What does C do?** C controls the trade-off between maximizing the margin and minimizing classification error:

Higher C (like 10) → Less tolerance for misclassification → tries to fit the training data more strictly.

Lower C (like 0.1) → More tolerant of misclassifications → allows a wider margin.

```
In [21]: model_1 = SVC(C=10)  
  
model_1.fit(X_train, y_train)
```


Out [21]:

```
▼ SVC  
SVC(C=10)
```

```
In [22]: model_1.score(X_test, y_test)
```

Out [22]: 0.9333333333333333

Training SVM with gamma=10 (Kernel Parameter) .

 **What does gamma do?** gamma controls the influence of a single training point in the decision boundary:

Higher gamma → More influence from each individual training point → creates a more complex decision boundary (can lead to overfitting).

Lower gamma → Each point has less influence → creates a simpler, smoother decision boundary (can lead to underfitting).

```
In [23]: model_2 = SVC(gamma=10)  
  
model_2.fit(X_train, y_train)
```

Out [23]:

▼ SVC
SVC(gamma=10)

In [24]: `model_2.score(X_test, y_test)`

Out [24]: 0.9

Training SVM with a Linear Kernel

📌 What does kernel='linear' mean? The linear kernel assumes that the data is linearly separable, meaning that a straight line (or hyperplane) can divide the classes in the feature space.

This is effective when the data is already well-separated or nearly linear.

```
In [25]: model_3 = SVC(kernel='linear')
model_3.fit(X_train, y_train)
```

Out [25]:

▼ SVC
SVC(kernel='linear')

In [26]: `model_3.score(X_test, y_test)`

Out [26]: 0.9333333333333333

📌 **What is Logistic Regression?** Logistic Regression is a probabilistic model used to predict class probabilities and is often used for binary classification. The model outputs a probability score between 0 and 1, representing the likelihood of belonging to a class.

In [27]: `model = LogisticRegression()`In [28]: `model.fit(X_train, y_train)`

```
/opt/conda/lib/python3.10/site-packages/sklearn/linear_model/_logistic.py:458:
ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

```
Increase the number of iterations (max_iter) or scale the data as shown in:
https://scikit-learn.org/stable/modules/preprocessing.html
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
n_iter_i = _check_optimize_result(
```

Out [28]:

▼ LogisticRegression
LogisticRegression()

```
In [29]: model.score(X_test, y_test)
```

```
Out[29]: 0.9333333333333333
```

📌 **What is a Confusion Matrix?** A confusion matrix is a table used to evaluate the performance of a classification model. It provides the following:

True Positives (TP): Correctly predicted positive samples.

True Negatives (TN): Correctly predicted negative samples.

False Positives (FP): Incorrectly predicted as positive.

False Negatives (FN): Incorrectly predicted as negative.

For multi-class classification (like the iris dataset), it will be a square matrix where:

Rows represent the true class.

Columns represent the predicted class.


```
In [30]: y_pred = model.predict(X_test)
confusion_m = confusion_matrix(y_test, y_pred)
confusion_m
```

```
Out[30]: array([[11,  0,  0],
               [ 0,  9,  2],
               [ 0,  0,  8]])
```

📌 **Why use a Heatmap?** A heatmap is a great way to visualize the confusion matrix because it makes it easy to see which classes are most frequently misclassified. The distribution of errors across classes.

```
In [31]: plt.figure(figsize=(7,5))
sns.heatmap(confusion_m, annot=True)
plt.xlabel('predicted')
plt.ylabel('Truth')
```

```
Out[31]: Text(58.22222222222214, 0.5, 'Truth')
```

📌 **Why use a Classification Report?** The classification report provides a more comprehensive view of your model's performance across each class: Precision is important when false positives are costly. Recall is important when false negatives are costly. F1-score balances precision and recall, which is useful when you need a single metric for model performance.

```
In [32]: print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	11
1	1.00	0.82	0.90	11
2	0.80	1.00	0.89	8
accuracy			0.93	30
macro avg	0.93	0.94	0.93	30
weighted avg	0.95	0.93	0.93	30

🇮🇹 **Classification Report** The classification report shows key metrics for each class (0, 1, and 2), as well as overall accuracy and averages.

Class-wise Metrics: Class 0 (Setosa):

Precision: 1.00 – Every prediction of class 0 is correct.

Recall: 1.00 – All actual class 0 instances were correctly predicted.

F1-score: 1.00 – Perfect balance between precision and recall.

Support: 12 – There are 12 instances of class 0 in the test set.

Class 1 (Versicolor):

Precision: 1.00 – All predictions for class 1 are correct.

Recall: 0.91 – 91% of the actual class 1 instances were correctly predicted.

F1-score: 0.95 – High balance between precision and recall.

Support: 11 – There are 11 instances of class 1.

Class 2 (Virginica):

Precision: 0.88 – 88% of predicted class 2 instances are correct.

Recall: 1.00 – All actual class 2 instances were correctly predicted.

F1-score: 0.93 – Good balance between precision and recall.

Support: 7 – There are 7 instances of class 2.

```
In [33]: y_test.shape
```

```
Out[33]: (30,)
```

Accuracy

```
In [34]: round((12+10+7)/(12+10+1+7),2)
```

```
Out[34]: 0.97
```

Precision for 1 class

```
In [35]: round(10/(10+0),2)
```

```
Out[35]: 1.0
```

```
In [36]: round(7/(7+1),2)
```

```
Out[36]: 0.88
```

Recall mean total Truth

```
In [37]: round(10/(10+1),2)
```

```
Out[37]: 0.91
```

```
In [38]: round(12/(12+0),2)
```

Out[38]: 1.0

F1-Score

$2((\text{precision} \cdot \text{recall}) / (\text{precision} + \text{recall}))$


In [39]: `round(2*(1*0.91)/(1+0.91),2)`

Out[39]: 0.95

In [40]: `round(2*(0.88*1)/(0.88+1),2)`

Out[40]: 0.94

.

 **Interpretation:** The model has performed excellently with an overall accuracy of 97%. Class 1 has a slightly lower recall (0.91), meaning some instances of Versicolor were misclassified. The F1-scores are close to perfect for all classes, indicating a good balance between precision and recall.

In []: